Glossary {-}

The terms this book leans on, defined once. Some have ancestors, credited in the chapters that introduce them. The definitions here are how this book uses them.

The Alexandria Problem — Not an event but a process: the library doesn't burn in one dramatic fire, it burns one skipped conversation at a time. The knowledge that used to accumulate through proximity and repetition stops the moment AI removes the friction of asking. Named for what happens when nobody notices the smoke until it's already gone.

Constructive friction — Disagreement from someone independent of your framing, arriving before the agent runs. The mechanism that tests shared understanding, and the thing an ever-agreeing tool cannot supply. Not an obstacle to remove; the reason the team's judgment is worth more than the model's agreement.

Direction — The running answer to "where does this need to go," held closely enough that the next decision can be checked against it: whether today's choice still makes sense six months from now, whether the system is becoming one thing or three. Judgment, once it's shared.

Epistemic safety — The conditions where people reveal what they don't know. Sharper than psychological safety: the junior asks without calculating the risk first, the senior shares the reasoning and not just the conclusion, the half-formed thought gets voiced. What the room runs on.

The gate — The pass condition that ends a Spec Session: intent stated, non-goals listed, edge cases answered or explicitly deferred, no sentence a new team member would misread. Also the unit of commitment: the team commits to specs through the gate, not to estimated capacity.

Left of the loop — The agent loop has a right side, where the agent runs: implementation, review, iteration. And it has a left side, where the thinking happens: what to build, why, and what done looks like. Left of the loop is the human work, the shared understanding built while the work is still on the page. This book is an argument for that side.

Minimum viable context — As little as possible, as much as needed. The right context, not just enough of it: what's assumed, what's load-bearing, what would change the answer if it were different. Includes surfacing the actual problem before asking about the assumed solution.

Predictability — Whether the team can tell a stakeholder when something will be done, and then do it. Not velocity, not throughput, not raw output: the metric that can't be gamed by one person and an agent for an afternoon. A team metric, not an individual one, because it takes shared understanding to earn.

Recognition work — Catching the wrong shape before it compounds: the drift from intent nobody had encoded yet, before there's a test to write. Quiet, constant, and learned from being in the system: the check an eval can't run.

The room — Wherever a team builds shared understanding together. A condition, not a location; the best room can have no walls at all. It has to be made, and protected, because everything the agent produces is downstream of what happens in it.

The spec — Shared understanding of what to build and why, held in common. Not a document; the document is the record of the convergence. Where the formal sense is meant — a contract, like the OpenFeature specification — this book uses the full word.

Spec Session — The deliberate moment where a team converges on a spec before the agent runs: intent challenged, constraints named, edge cases answered, consent recorded. It inherits refinement's content and mob programming's mechanics, and it replaces the walk to the desk. A working template waits at the back.

Stakeholder Navigator — The role that owns the organizational interface after product thinking distributes into the team: carries the spec to whoever the outcome touches, manages the review cycle, and protects the room from the noise while the team thinks. The part of the old Product Owner that doesn't distribute.

Weighted validation — Agreement from someone independent of your framing, with their own hard-won experience and their own reasons to disagree, who could have said no. It changes something an AI's agreement can't, because it costs the person something to give. One of four things only another person can provide, before the agent runs.

The XY problem — Someone has problem X, thinks the solution is Y, and asks for help with Y without saying what X was. The answer to Y lands and solves nothing, because the person helping never found out about X. Coined by Mark Jason Dominus; minimum viable context is what prevents it.